

# **MINERÍA DE DATOS DESCRIPTIVA**

# **LIBRERIA PANDAS DE PYTHON**

# Pandas: Panel Data System

- Python library to provide data analysis features
- Built on top of NumPy (SciPy, matplotlib)
- Key components provided by pandas:
  - Series
  - DataFrame
- The ideal tool for data scientists
  - Munging data
  - Cleaning data
  - Analyzing
  - Modeling
  - Organizing the result of the analysis into a suitable form for plotting or tabular display

# Pandas: Axis Indexing

- Every axis has an index
- Highly optimized structure
- Hierarchical indexing
  - Semantics: a tuple at each tick
  - Enables easy group by selection
- Group by and join-type operations



A	1
	2
	3
B	1
	2
	3
	4

# Pandas: Series

- One-dimensional array-like object containing data and labels
  - Sublass of numpy.ndarray
  - Data: can be of any type
  - Index labels need not be ordered
  - Duplicates are possible at the cost of a reduced functionality

index	values
A	→ 5
B	→ 6
C	→ 12
D	→ -5
E	→ 6.7

# Pandas: Series construction

- Lot of ways to build Series
  - From lists
  - From dictionaries

```
In [35]: import pandas as pd
In [36]: s = pd.Series(list('abcdef'))
In [37]: s
Out[37]:
0    a
1    b
2    c
3    d
4    e
5    f
dtype: object
In [38]: s = pd.Series([2,4,6,8])
In [39]: s
Out[39]:
0    2
1    4
2    6
3    8
dtype: int64
```

```
In [45]: d = dict({'a':-0.889, 'b':-2.187,'c':1.102})
In [46]: d
Out[46]: {'a': -0.889, 'b': -2.187, 'c': 1.102}
In [47]: pd.Series(d)
Out[47]:
a    -0.889
b    -2.187
c    1.102
dtype: float64
```

# Pandas: Series indexing

- A Series index can be specified
  - Single values can be selected by index
  - Multiple values can be selected with multiple indexes

```
In [64]: s = pd.Series([2,4,6,8], index = ['f','a','c','e'])

In [65]: s
Out[65]:
f    2
a    4
c    6
e    8
dtype: int64

In [66]: s['a']
Out[66]: 4

In [67]: s[['a','c']]
Out[67]:
a    4
c    6
dtype: int64
```

# Pandas: Series indexing

- Think of Series as a fixed-length, ordered dictionary
  - However index items do not have to be unique

```
In [68]: s = pd.Series(range(4), index = list('abab'))  
In [69]: s  
Out[69]:  
a    0  
b    1  
a    2  
b    3  
dtype: int64  
  
In [70]: s['a']  
Out[70]:  
a    0  
a    2  
dtype: int64  
  
In [71]: s['a'][0]  
Out[71]: 0
```

# Pandas: Series operations

- Series works with NumPy
  - Filtering
  - Mathematical operations

```
In [90]: s = pd.Series(range(4), index = list('asdf'))  
  
In [91]: s  
Out[91]:  
a    0  
s    1  
d    2  
f    3  
dtype: int64  
  
In [92]: np.sin(s)  
Out[92]:  
a    0.000000  
s    0.841471  
d    0.909297  
f    0.141120  
dtype: float64
```

```
In [95]: s  
Out[95]:  
a    0  
s    1  
d    2  
f    3  
dtype: int64  
  
In [96]: s[s>=2]  
Out[96]:  
d    2  
f    3  
dtype: int64  
  
In [97]: s*2  
Out[97]:  
a    0  
s    2  
d    4  
f    6  
dtype: int64
```

# Pandas: Series (missing data)

- Series can accommodate incomplete data
- Unlike in a NumPy ndarray, data is automatically aligned
  - ▣ Binary operations are joins

$$\begin{array}{|c|c|} \hline \textbf{B} & 1 \\ \hline \textbf{C} & 2 \\ \hline \textbf{D} & 3 \\ \hline \textbf{E} & 4 \\ \hline \end{array} + \begin{array}{|c|c|} \hline \textbf{A} & 0 \\ \hline \textbf{B} & 1 \\ \hline \textbf{C} & 2 \\ \hline \textbf{D} & 3 \\ \hline \end{array} = \begin{array}{|c|c|} \hline \textbf{A} & \textbf{NA} \\ \hline \textbf{B} & 2 \\ \hline \textbf{C} & 4 \\ \hline \textbf{D} & 6 \\ \hline \textbf{E} & \textbf{NA} \\ \hline \end{array}$$

```
In [98]: d
Out[98]: {'a': -0.889, 'b': -2.187, 'c': 1.102}

In [99]: s = pd.Series(d)

In [100]: s
Out[100]:
a    -0.889
b    -2.187
c    1.102
dtype: float64

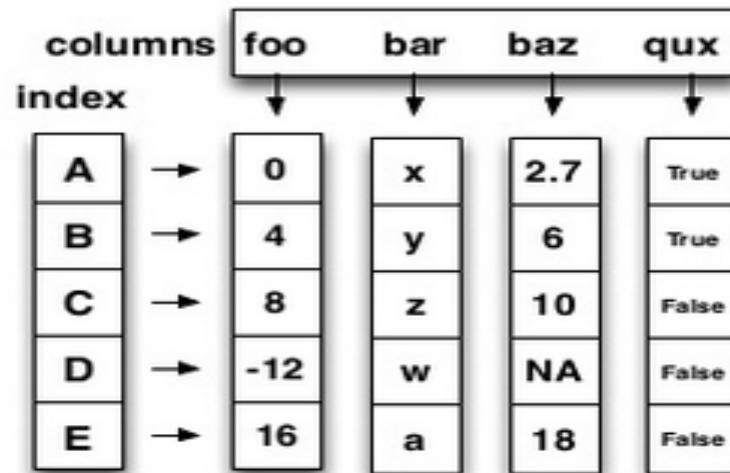
In [101]: s = pd.Series(d, list('abcd'))

In [102]: s
Out[102]:
a    -0.889
b    -2.187
c    1.102
d    NaN
dtype: float64

In [103]: s*2
Out[103]:
a    -1.778
b    -4.374
c    2.204
d    NaN
dtype: float64
```

# Pandas: DataFrames

- Spreadsheet-like data structure containing an ordered collection of columns
  - Data that can be represented as tables: Rows and columns (with their own indexes)
    - Each row is a different object
    - Columns represent attributes of the objects
  - Each column can have a different type
    - But a column is homogenously typed
- Consider as a dictionary of Series (with shared index)
- Size mutable: insert and delete columns



# Pandas: DataFrames generation

- DataFrame generation from a dictionary of equal-length lists
- Generation from a dictionary of dictionaries

```
In [112]: s = pd.Series([1,2,3])
In [113]: s
Out[113]:
0    1
1    2
2    3
dtype: int64
In [114]: d1 = {'one':s+s, 'two':s*s}
In [115]: d1
Out[115]:
{'one': 0    2
 1    4
 2    6
  dtype: int64, 'two': 0    1
 1    4
 2    9
  dtype: int64}
In [116]: frame1 = pd.DataFrame(d1)
In [117]: frame1
Out[117]:
   one  two
0    2    1
1    4    4
2    6    9
```

```
In [107]: data = {'state': ['FL','FL','GA','GA','GA'],'year':[2010,2011,2008,2010,2011], 'pop': [18.8,19.1,9.7,9.7,9.8]}
In [108]: frame = pd.DataFrame(data)
In [109]: frame
Out[109]:
   pop state  year
0  18.8    FL  2010
1  19.1    FL  2011
2   9.7    GA  2008
3   9.7    GA  2010
4   9.8    GA  2011
```

```
In [118]: dd = {'FL': {2010:18.8, 2011: 19.1}, 'GA': {2008: 9.7, 2010: 9.7, 2011: 9.8}}
In [119]: frameDD = pd.DataFrame(dd)
In [120]: frameDD
Out[120]:
          FL    GA
2008    NaN  9.7
2010  18.8  9.7
2011  19.1  9.8
```

# Pandas: Data loading

- Pandas supports several ways to handle **data loading**
- Text file data
  - **read\_csv**
    - filepath\_or\_buffer: the file to read (it can be a URL)
    - sep : Delimiter to use (default ',')
    - header: number of row where the column names are specified (default 0)
      - If there are no column names set it to None
    - names: list of the column names (header set to 0 to change them)
    - nrows : Number of rows of file to read. Useful for reading pieces of large files (default None: all file)
    - More information at this [website](#)
  - **to\_csv**
    - filepath: name of the file('fileName.csv')
      - df.to\_csv('example.csv'): writes the content of the DataFrame df in the file example.csv
  - Structured data: JSON, XML, HTML
  - Specific libraries
- Excel and database

# Pandas: DataFrames slicing

- Select column by name
- To select data we can use the methods:
  - *loc* allows one to select data using row and column labels
  - *iloc* allows one to select data using row and column position
  - *ix* uses both labels and positions
- Using DataFrames they can return
  - A Series if only the row position is specified
  - A single value if the row and column are specified
    - In the case of a Series they return a single value

```
In [84]: df = pd.DataFrame({'uno' : pd.Series([1, 2, 3], index=['a', 'b', 'c']),  
...: 'dos' : pd.Series([4,5, 6,7], index=['a', 'b', 'c', 'd'])})  
  
In [85]: df  
Out[85]:  
      dos  uno  
a      4    1  
b      5    2  
c      6    3  
d      7   NaN  
  
In [86]: df.loc['a']  
Out[86]:  
dos    4  
uno    1  
Name: a, dtype: float64  
  
In [87]: df.loc['a', 'uno']  
Out[87]: 1.0  
  
In [88]: df.iloc[2]  
Out[88]:  
dos    6  
uno    3  
Name: c, dtype: float64  
  
In [89]: df.iloc[2,0]  
Out[89]: 6.0  
  
In [90]: df.ix[0]  
Out[90]:  
dos    4  
uno    1  
Name: a, dtype: float64  
  
In [91]: df.ix['b',0]  
Out[91]: 5.0  
  
In [121]: frame['state']  
Out[121]:  
0    FL  
1    FL  
2    GA  
3    GA  
4    GA  
Name: state, dtype: object
```

# Pandas: DataFrames indexing

## □ Boolean indexing

```
In [136]: frameDD
Out[136]:
      FL    GA
2008   NaN  9.7
2010  18.8  9.7
2011  19.1  9.8

In [137]: frameDD < 9.8
Out[137]:
      FL    GA
2008  False  True
2010  False  True
2011  False False

In [138]: frameDD[frameDD < 9.8] = 0

In [139]: frameDD
Out[139]:
      FL    GA
2008   NaN  0.0
2010  18.8  0.0
2011  19.1  9.8
```

```
In [171]: f
Out[171]:
      one    two
0      2      2
1      4      4
2      6      6

In [172]: aux = f['one'] > 3

In [173]: f[aux]
Out[173]:
      one    two
1      4      4
2      6      6
```

# Pandas: DataFrames indexing

- You can obtain all the indexes using the `index` attribute
- The index can be changed: `set_index` function
  - `set_index` returns a new DataFrame
  - If you want to modify the actual one you have to specify it `setting` the `inplace` parameter to `True`
- The `index` can be reset to the original one: `reset_index`

```
print(users.set_index('user_id').head())
print('\n')

print(users.head())
print("\n^^^ I didn't actually change the DataFrame. ^^^\n")
```

	user_id	age	sex	occupation	zip_code
1	1	24	M	technician	85711
2	2	53	F	other	94043
3	3	23	M	writer	32067
4	4	24	M	technician	43537
5	5	33	F	other	15213

```
users.set_index('user_id', inplace=True)
users.head()
```

	user_id	age	sex	occupation	zip_code
1	1	24	M	technician	85711
2	2	53	F	other	94043
3	3	23	M	writer	32067
4	4	24	M	technician	43537
5	5	33	F	other	15213

# Pandas: DataFrames iterating

- *iteritems* and *iterrows* allow one to iterate over a DataFrame
  - Iteritems returns the DataFrame by column
  - Iterrows returns the DataFrame by rows

```
In [96]: df
Out[96]:
      dos  uno
a    4    1
b    5    2
c    6    3
d    7  NaN

In [97]: for column, values in df.iteritems():
...:     print column
...:     print list(values)
...:
dos
[4, 5, 6, 7]
uno
[1.0, 2.0, 3.0, nan]

In [98]: for row, values in df.iterrows():
...:     print row
...:     print list(values)
...:
a
[4.0, 1.0]
b
[5.0, 2.0]
c
[6.0, 3.0]
d
[7.0, nan]
```

# Pandas: DataFrames visualization

- The function `head()` displays the first five records of the dataset
  - The function `tail()` displays the last five
  - A different number of records (n) can be specified by parameter
    - `head(n)` or `tail(n)`
- In both cases you can visualize the data of only one field or a set of fields
  - `df[['field']].head()` or `df[['field','field1']].head()`
- You can also make a selection of the fields using conditions
  - `df[df.'field' > X].head()`

# Pandas: DataFrames resizing

- Column addition

- ▣ By computation

```
In [146]: frame['calc'] = frame['pop']*2
In [147]: frame
Out[147]:
   pop state  year  other  calc
0  18.8    FL  2010     100  37.6
1  19.1    FL  2011     100  38.2
2  9.7     GA  2008     100  19.4
3  9.7     GA  2010     100  19.4
4  9.8     GA  2011     100  19.6
```

- ▣ By direct assignment

```
In [144]: frame['other'] = 100
In [145]: frame
Out[145]:
   pop state  year  other
0  18.8    FL  2010    100
1  19.1    FL  2011    100
2  9.7     GA  2008    100
3  9.7     GA  2010    100
4  9.8     GA  2011    100
```

# Pandas: DataFrames reshaping

- Pivot tables: you can [create a table](#) having the result of an aggregation over specific fields
  - `data.pivot_table(values=None, index=None, columns=None, aggfunc='mean', fill_value=None, margins=False, dropna=True, margins_name='All')`
    - Data: a DataFrame
    - Values: the column used to perform the aggregation
    - Index: the field(s) used as index in the pivot table
    - Columns: the field(s) to use as columns in the pivot table
      - Index and columns will be used to group the values and perform the aggregation over them
    - Aggfunc: the aggregation to apply
    - Fill\_value: the value used to fill empty fields

```
>>> df
      A   B   C   D
0  foo one small  1
1  foo one large  2
2  foo one large  2
3  foo two small  3
4  foo two small  3
5  bar one large  4
6  bar one small  5
7  bar two small  6
8  bar two large  7
```

```
>>> table = df.pivot_table(values='D', index=['A', 'B'],
...                           columns=['C'], aggfunc=np.sum)
...
>>> table
          small  large
foo one    1     4
              two   6   NaN
bar one    5     4
              two   6     7
```

# Pandas: DataFrames operations

## □ Descriptive statistics

Function	Description
count	Number of non-null observations
sum	Sum of values
mean	Mean of values
mad	Mean absolute deviation
median	Arithmetic median of values
min	Minimum
max	Maximum
mode	Mode
abs	Absolute Value
prod	Product of values
std	Unbiased standard deviation
var	Unbiased variance
sem	Unbiased standard error of the mean
skew	Unbiased skewness (3rd moment)
kurt	Unbiased kurtosis (4th moment)
quantile	Sample quantile (value at %)
cumsum	Cumulative sum
cumprod	Cumulative product
cummax	Cumulative maximum
cummin	Cumulative minimum

```
In [148]: frameDD
Out[148]:
      FL   GA
2008  NaN  0.0
2010  18.8  0.0
2011  19.1  9.8

In [149]: frameDD.sum()
Out[149]:
FL    37.9
GA     9.8
dtype: float64

In [150]: frameDD.mean()
Out[150]:
FL    18.950000
GA     3.266667
dtype: float64

In [151]: frameDD.describe()
Out[151]:
              FL         GA
count    2.000000  3.000000
mean    18.950000  3.266667
std     0.212132  5.658033
min    18.800000  0.000000
25%    18.875000  0.000000
50%    18.950000  0.000000
75%    19.025000  4.900000
max    19.100000  9.800000
```

# Pandas: DataFrames operations

- More functions
  - [Size](#): it counts the total number of values (not by columns)
  - [Sort\\_values](#): to sort a Series
    - If you have a DataFrame you have to specify the field name to order by
      - It is specified as a tuple
    - The parameter named ascending allows to sort from the largest to the less number (True, by default) or viceversa (False)

# Pandas: DataFrames operations

- The function `apply` allows one to apply a function over the columns or rows
- The function `applymap` allows one to apply a function element wise

```
In [110]: df = pd.DataFrame({'uno' : pd.Series([1, 2, 3], index=['a', 'b', 'c']),
...: 'dos' : pd.Series([4,5, 6,7], index=['a', 'b', 'c', 'd']), 'tres' : pd.Series([8,9,10,11], index=['a',
'b', 'c', 'd'])})
```

```
In [111]: df
Out[111]:
   dos  tres  uno
a    4     8   1
b    5     9   2
c    6    10   3
d    7    11  NaN
```

```
In [112]: df.apply(lambda x: x.max())
Out[112]:
dos      7
tres     11
uno      3
dtype: float64
```

```
In [113]: df.apply(lambda x: [x.max(), x.min()])
Out[113]:
dos      [7.0, 4.0]
tres     [11.0, 8.0]
uno      [3.0, 1.0]
dtype: object
```

```
In [114]: df.apply(lambda x: [x.max(), x.min()],axis=0)
Out[114]:
dos      [7.0, 4.0]
tres     [11.0, 8.0]
uno      [3.0, 1.0]
dtype: object
```

```
In [115]: df.apply(lambda x: [x.max(), x.min()],axis=1)
Out[115]:
a      [8.0, 1.0]
b      [9.0, 2.0]
c      [10.0, 3.0]
d      [11.0, 7.0]
dtype: object
```

```
In [116]: df.applymap(lambda x: (x, 2**x))
Out[116]:
              dos          tres          uno
a  (4, 16)    (8, 256)  (1.0, 2.0)
b  (5, 32)    (9, 512)  (2.0, 4.0)
c  (6, 64)    (10, 1024) (3.0, 8.0)
d  (7, 128)   (11, 2048) (nan, nan)
```

```
In [117]: df.applymap(lambda x: [x.max(), x.min()])
Out[117]:
              dos          tres          uno
a  [4, 4]     [8, 8]  [1.0, 1.0]
b  [5, 5]     [9, 9]  [2.0, 2.0]
c  [6, 6]     [10, 10] [3.0, 3.0]
d  [7, 7]     [11, 11] [nan, nan]
```

# DataFrames joining

- Like SQL's JOIN clause, `pandas.merge` allows two DataFrames to be joined on one or more keys
  - To specify the field to join by use the `parameter on`
- By default, `pandas.merge` operates as an `inner join`, which can be changed using the `how parameter`
  - `how : {'left', 'right', 'outer', 'inner'}`, default 'inner'
    - left: use only keys from left frame (SQL: left outer join)
    - right: use only keys from right frame (SQL: right outer join)
    - outer: use union of keys from both frames (SQL: full outer join)
    - inner: use intersection of keys from both frames (SQL: inner join)

# DataFrames joining example

```
left_frame = pd.DataFrame({'key': range(5),
                           'left_value': ['a', 'b', 'c', 'd', 'e']})
right_frame = pd.DataFrame({'key': range(2, 7),
                            'right_value': ['f', 'g', 'h', 'i', 'j']})
```

```
key left_value
0   0      a
1   1      b
2   2      c
3   3      d
4   4      e
```

inner join (default)

```
pd.merge(left_frame, right_frame, on='key', how='inner')
```

	key	left_value	right_value
0	2	c	f
1	3	d	g
2	4	e	h

```
key right_value
0   2      f
1   3      g
2   4      h
3   5      i
4   6      j
```

right outer join

```
pd.merge(left_frame, right_frame, on='key', how='right')
```

	key	left_value	right_value
0	2	c	f
1	3	d	g
2	4	e	h
3	5	NaN	i
4	6	NaN	j

left outer join

```
pd.merge(left_frame, right_frame, on='key', how='left')
```

	key	left_value	right_value
0	0	a	NaN
1	1	b	NaN
2	2	c	f
3	3	d	g
4	4	e	h

full outer join

```
pd.merge(left_frame, right_frame, on='key', how='outer')
```

	key	left_value	right_value
0	0	a	NaN
1	1	b	NaN
2	2	c	f
3	3	d	g
4	4	e	h
5	5	NaN	i
6	6	NaN	j

# DataFrames combining

- ❑ Pandas also provides a way to combine DataFrames along an axis - `pandas.concat`
  - ❑ The function is equivalent to SQL's UNION clause

```
key left_value
0   0     a
1   1     b
2   2     c
3   3     d
4   4     e

key right_value
0   2     f
1   3     g
2   4     h
3   5     i
4   6     j
```

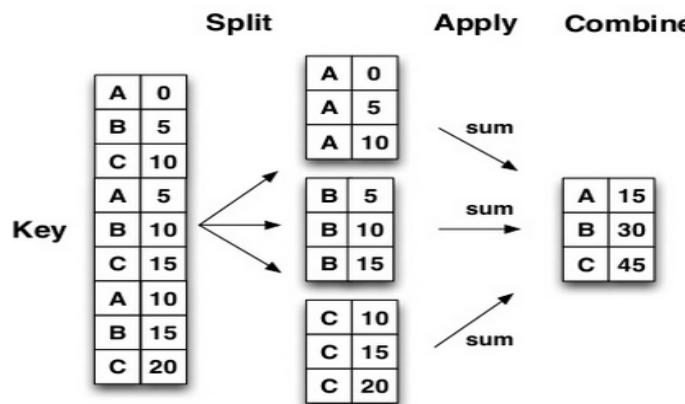
	key	left_value	right_value
0	0	a	NaN
1	1	b	NaN
2	2	c	NaN
3	3	d	NaN
4	4	e	NaN
	0	NaN	f
	1	NaN	g
	2	NaN	h
	3	NaN	i
	4	NaN	j

```
pd.concat([left_frame, right_frame], axis=1)
```

	key	left_value	key	right_value
0	0	a	2	f
1	1	b	3	g
2	2	c	4	h
3	3	d	5	i
4	4	e	6	j

# Pandas: DataFrame aggregation

- Groupby essentially splits the data into different groups depending on a variable of your choice
  - The groupby() function returns a GroupBy object, which have an attribute named `groups` that is a dictionary containing
    - Keys: computed unique groups
    - Values: the axis labels belonging to each group
    - These values can be accessed by the functions `keys()`, `values()` and `items()`
  - Functions like `max()`, `min()`, `mean()`, `first()`, `last()` can be quickly applied to the GroupBy object to obtain summary statistics for each group



# Pandas: DataFrame aggregation example

- Dataset contains 830 entries from my mobile phone log spanning a total time of 5 months

index	date	duration	item	month	network	network_type
0	15/10/14 06:58	34.429	data	2014-11	data	data
1	15/10/14 06:58	13.000	call	2014-11	Vodafone	mobile
2	15/10/14 14:46	23.000	call	2014-11	Meteor	mobile
3	15/10/14 14:48	4.000	call	2014-11	Tesco	mobile
4	15/10/14 17:27	4.000	call	2014-11	Tesco	mobile
5	15/10/14 18:55	4.000	call	2014-11	Tesco	mobile
6	16/10/14 06:58	34.429	data	2014-11	data	data
7	16/10/14 15:01	602.000	call	2014-11	Three	mobile
8	16/10/14 15:12	1050.000	call	2014-11	Three	mobile
9	16/10/14 15:30	19.000	call	2014-11	voicemail	voicemail
10	16/10/14 16:21	1183.000	call	2014-11	Three	mobile

## Simple statistics from Pandas

```
1 # How many rows the dataset
2 data['item'].count()
3 Out[38]: 830
4
5 # What was the longest phone call / data entry?
6 data['duration'].max()
7 Out[39]: 10528.0
8
9 # How many seconds of phone calls are recorded in total?
10 data['duration'][data['item'] == 'call'].sum()
11 Out[40]: 92321.0
12
13 # How many entries are there for each month?
14 data['month'].value_counts()
15 Out[41]:
16 2014-11    230
17 2015-01    205
18 2014-12    157
19 2015-02    137
20 2015-03    101
```

```
Pandas Groupby Functionality
1 # Get the first entry for each month
2 data.groupby('month').first()
3 Out[69]:
4
5          date   duration item network network_type
6 2014-11 2014-10-15 06:58:00 34.429 data     data     data
7 2014-12 2014-11-13 06:58:00 34.429 data     data     data
8 2015-01 2014-12-13 06:58:00 34.429 data     data     data
9 2015-02 2015-01-13 06:58:00 34.429 data     data     data
10 2015-03 2015-02-12 20:15:00 69.000 call landline landline
11
12 # Get the sum of the durations per month
13 data.groupby('month')['duration'].sum()
14 Out[70]:
15
16 2014-11    26639.441
17 2014-12    14641.870
18 2015-01    18223.299
19 2015-02    15522.299
20 2015-03    22750.441
21 Name: duration, dtype: float64
22
23 # Get the number of dates / entries in each month
24 data.groupby('month')['date'].count()
25 Out[74]:
26
27 2014-11    230
28 2014-12    157
29 2015-01    205
30 2015-02    137
31 2015-03    101
32 Name: date, dtype: int64
33
34 # What is the sum of durations, for calls only, to each network
35 data[data['item'] == 'call'].groupby('network')['duration'].sum()
36 Out[78]:
37
38 Meteor 7200
39 Tesco 13828
40 Three 36464
41 Vodafone 14621
42 landline 18433
43 voicemail 1775
```

# Pandas: DataFrame aggregation example

- Data can be grouped by more than one variable

```
Grouping by multiple variables
1 # How many calls, sms, and data entries are in each month?
2 data.groupby(['month', 'item'])['date'].count()
3 Out[76]:
4    month      item
5 2014-11    call    107
6          data    29
7          sms    94
8 2014-12    call    79
9          data    30
10         sms    48
11 2015-01    call    88
12         data    31
13         sms    86
14 2015-02    call    67
15         data    31
16         sms    39
17 2015-03    call    47
18         data    29
19         sms    25
20 Name: date, dtype: int64
21
22 # How many calls, texts, and data are sent per month, split by network_type?
23 data.groupby(['month', 'network_type'])['date'].count()
24 Out[82]:
25    month network_type
26 2014-11   data  29
27        landline  5
28        mobile 189
29        special  1
30    voicemail  6
31 2014-12   data  30
32        landline  7
33        mobile 108
```

# Pandas: DataFrame aggregation

## □ Multiple Statistics per Group: `agg()` function

```
Work on this column  
aggregations = {  
    'duration': {  
        'total_duration': 'sum',  
        'average_duration': 'mean',  
        'num_calls': 'count'  
    },  
    'date': {  
        'max_date': 'max',  
        'min_date': 'min',  
        'num_days': lambda x: max(x) - min(x)  
    },  
    'network': ["count", "max"]  
}  
data[data['item'] == 'call'].groupby('month').agg(aggregations)
```

Name the results  
Perform these operations

Out[47]:

	duration			date				network	
	average_duration	num_calls	total_duration	max_date	num_days	min_date	count	max	
month									
2014-11	238.757009	107	25547	2014-11-12 19:01:00	28 days 12:03:00	2014-10-15 06:58:00	107	voicemail	
2014-12	171.658228	79	13561	2014-12-14 19:54:00	30 days 02:30:00	2014-11-14 17:24:00	79	voicemail	
2015-01	193.977273	88	17070	2015-01-14 20:47:00	30 days 00:44:00	2014-12-15 20:03:00	88	voicemail	
2015-02	215.164179	67	14416	2015-02-09 17:54:00	25 days 07:18:00	2015-01-15 10:36:00	67	voicemail	
2015-03	462.276596	47	21727	2015-03-04 12:29:00	19 days 16:14:00	2015-02-12 20:15:00	47	voicemail	

# Pandas: DataFrame aggregation

- The agg() function returns MultiIndexes
  - In the previous example
    - ('duration', 'total\_duration'), ('duration', 'average\_duration'), ('duration', 'num\_calls')
    - ('date', 'max\_date'), ('date', 'min\_date'), ('date', 'num\_days')
    - ('network', 'count'), ('network', 'max')
- If you want to operate over the result of the agg function
  - You have to specify the column by the corresponding MultiIndex
    - Example: `data.sort_values(('network', 'count'), ascending=False)`

# Pandas: visualization

- Pandas **visualization** uses the `matplotlib` package
  - `import matplotlib.pyplot as plt`
- The **plot method** on Series and DataFrame is just a simple wrapper around `plt.plot()`
- Basic plotting: **plot()**

```
In [2]: ts = pd.Series(np.random.randn(1000),  
index=pd.date_range('1/1/2000', periods=1000))
```

```
In [3]: ts = ts.cumsum()
```

```
In [4]: ts.plot()
```

```
In [5]: df =  
pd.DataFrame(np.random.randn(1000, 4),  
index=ts.index, columns=list('ABCD'))
```

```
In [6]: df = df.cumsum()
```

```
In [7]: plt.figure(); df.plot();
```

# Pandas: visualization

- Other options: DataFrame.plot.<kind>
  - df.plot.area
  - df.plot.bar (df.plot.bahr in horizontal)
  - df.plot.density
  - df.plot.hist
  - df.plot.line
  - df.plot.scatter
  - df.plot.box
  - df.plot.hexbin
  - df.plot.kde
  - df.plot.pie
- See the [documentation](#) for details

# More information about pandas

- <http://pandas.pydata.org/pandas-docs/version/0.18.0/pandas.pdf>